

S3 Versioning, AWS Glacier and Static Website Deployment on AWS S3

Due on August 30, 2021

Basanta Joshi, Ph.D.

Dipak Poudel

Ashlesh Pandey

PUL074BEX007

Contents

List of Figures	ii
List of Abbreviations	iii
1 S3 Versioning	1
1.1 Different scenarios in versioned bucket	2
1.1.1 PUT an object in versioning enabled bucket	2
1.1.2 DELETE an object from versioning enabled bucket	3
1.1.3 GET an object from versioning enabled bucket	3
1.1.4 Restoring a previous version of an object	5
2 AWS Glacier	6
3 Static Website Deployment on AWS S3	8

List of Figures

1	S3 versioning disabled versus enabled	1
2	PUT an object in versioning enabled bucket	2
3	DELETE an object from versioning enabled bucket	3
4	GET a delete marked object	4
5	GET a specific version of object by using ID	4
6	COPY a previous version of an object	5
7	Permanently DELETE the current version of an object	6
8	Log in to AWS console	9
9	Open S3 service from the AWS services list	10
10	Create a S3 bucket	11
11	Set an unique bucket name	12
12	Grant public access to bucket	13
13	Complete bucket creation	14
14	Select the created bucket	15
15	Upload objects to the bucket	16
16	Upload a file or a folder as needed	17
17	Select the file to upload	18
18	Upload selected file	19
19	Confirm the file upload	20
20	Make the object public	21
21	Review the selected object	22
22	Confirm public access	23
23	Copy the object url	24
24	Visit the url and access the website	25

List of Abbreviations

AWS Amazon Web Services.

CSS Cascading Style Sheets.

HTML Hyper Text Markup Language.

JS Javascript.

S3 Simple Storage Service.

1 S3 Versioning

Amazon Simple Storage Service (Amazon S3) is the object storage service provided by AWS. Objects are stored in a S3 bucket, whose versioning is disabled by default. S3 versioning, or simply versioning is a method that allows users of S3 buckets to keep multiple versions of an object in the same S3 bucket. It is particularly useful for,

1. Avoiding accidental deletion of an object as it creates a delete marker on the latest version of the deleted object essentially making the delete marker object the current version.
2. Avoiding unintentional overwriting of an object as it creates a new version of the object with the same key rather than overwriting the existing object.

A note-worthy point is that, S3 versioning is disabled by default, so if an user wishes to have versioning enabled, they need to do it either during bucket creation or at a later time.




Figure 1: S3 versioning disabled versus enabled

Objects placed in buckets with versioning disabled have their ID set to *null*. If a bucket contains objects prior to the enabling of the versioning, the objects will still have their IDs marked as *null*. Any new objects that are PUT in the bucket will get an auto-generated version ID that are composed of various unicode, UTF-8 encoded, URL-ready, opaque strings. To explain S3 versioning, simpler IDs are used in this report. This ID distinguishes the different versions of

the same key.

S3 buckets can exist in one of the following states,

1. Versioning disabled (default)
2. Versioning enabled
3. Versioning suspended

Once a bucket's versioning is enabled, it can never be disabled, rather can only be set to a versioning suspended state by an user with the required permissions.

1.1 Different scenarios in versioned bucket

1.1.1 PUT an object in versioning enabled bucket

When an user PUTs an object in a versioning enabled bucket with the key of an existing object, rather than overwriting the object, the new object is stored with a new version ID and becomes the latest version of that object. This way accidental overwriting is prevented.




Figure 2: PUT an object in versioning enabled bucket

1.1.2 DELETE an object from versioning enabled bucket

When an user DELETES an object from a versioning enabled bucket, rather than deleting the object, a delete marker is inserted with a new version ID and becomes the latest version of that object. This way accidental deletion is prevented as the older versions of the object are still stored in the same S3 bucket.

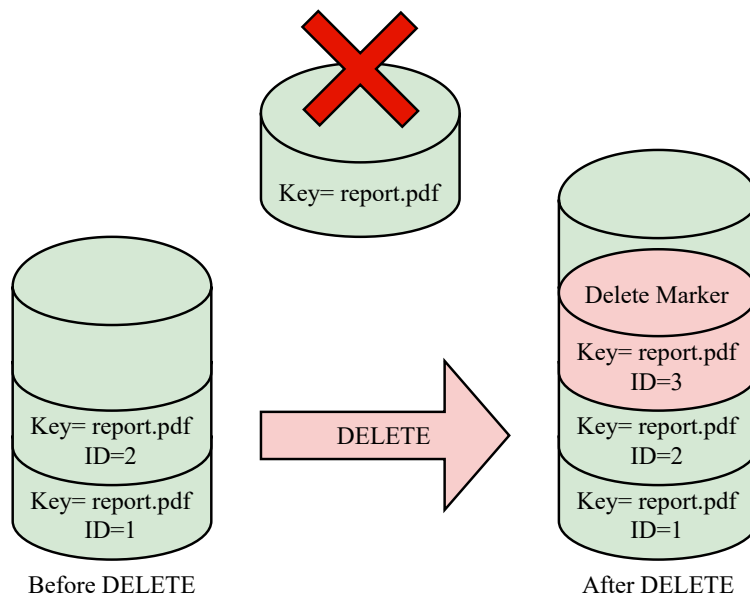


Figure 3: DELETE an object from versioning enabled bucket

1.1.3 GET an object from versioning enabled bucket

Since GET request returns the latest version of an object, a simple GET request to retrieve the report.pdf object in the current situation would return a '404 No Object Found' error since the latest version is a delete marker. This is because the delete marker identifies the object as deleted although there are older versions of the object existing in the same bucket. If the user wishes to retrieve an older version of the object that is not marked as deleted, then the GET request must be modified to point to the exact unique version ID that the user wants to get.

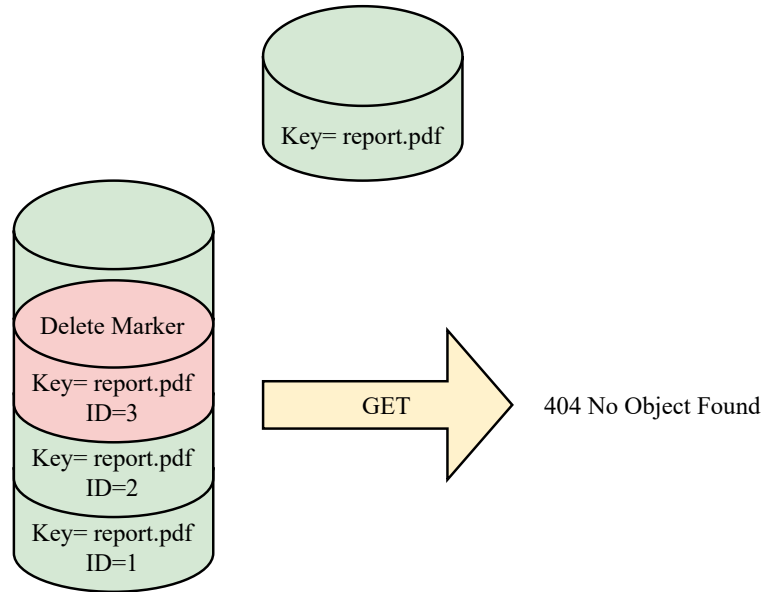


Figure 4: GET a delete marked object




Figure 5: GET a specific version of object by using ID

This helps in explaining how the different versions of the object are separately accessible with their unique version ID.

1.1.4 Restoring a previous version of an object

To restore previous versions of an object, one of two methods can be used,

1. Copy a pervious version of the object into the bucket.
2. Permanently delete the current version of the object.




Figure 6: COPY a previous version of an object

A copy of a previous version can be made into the same S3 bucket by first using GET request to retrieve the required ID followed by a PUT request. A note-worthy point during this method is that, the old version isn't moved and made the current version, rather a new object with a new unique ID is PUT into the bucket that is a copy of the requested version.




Figure 7: Permanently DELETE the current version of an object

Unlike Figure 3, where a general DELETE request inserts a delete marker as the current version, if the user specifies the ID of the version that they wish to delete, it is deleted permanently. This can be used to some extent to restore the immediately previous version of the object, but if there are multiple versions existing, and the need is to restore a version that is quite old, multiple versions of the object would need to be deleted, which might not be what the user wants.

2 AWS Glacier

AWS Glacier is a low-cost, secure and durable storage class under AWS S3 storage classes, mostly used for data archiving purposes. As the name suggests, S3 Glacier is used to store cold data, a term commonly used for data that is archived or infrequently accessed but preserved for future access. Many organizations have the need of storing cold data but surely won't want high-cost storage services to do so. To keep the cost of AWS Glacier as low as possible, AWS

provides data retrieval in one of three ways,

1. **Bulk retrieval:** To recover large amount of data in 5 to 12 hours.
2. **Standard retrieval:** To recover data in 3 to 5 hours.
3. **Expedited retrieval:** To recover data within 5 minutes.

Some use cases of AWS Glacier are,

1. **Digital preservation:** Many governmental firms and offices utilize Glacier service for digital preservation of the important data.
2. **Regulatory and compliance archiving:** Since health and financial firms require large amount of regulatory and compliance archives, they use Glacier services.
3. **Tape replacement:** With large upfront investment and maintenance needs, tape archiving can be replaced by Glacier archiving.
4. **Media assets storage:** Large media files that require a lot of space are stored in affordable Glacier buckets.

Some of the benefits of AWS Glacier are,

1. **Durability and scalability:** Since AWS Glacier is distributed across at least three availability zones, the durability of the storage is unquestioned. Similarly, organizations can scale their cloud archives as per their requirements.
2. **Lower cost:** AWS Glacier is one of the lowest-cost storage classes and also employs pay-as-you-go, hence reducing the cost tremendously.
3. **Security:** Glacier has support of various compliance standards and encryption that allow secure storage at low cost.

4. **Retrieval options:** As mentioned above, there are different options that a client can choose from while retrieving their data hence giving them choice based on retrieval time and cost.

3 Static Website Deployment on AWS S3

To deploy static websites on AWS S3, users need to create a storage bucket where the static code (HTML, CSS, JS) files are kept as objects. Necessary policies are then created for the public access of these static files and a S3 url for the HTML object makes the website accessible over the internet.

For this report, a simple HTML file is used.

```
1 <html>
2   <head>
3   </head>
4   <body>
5       <h1>Static website hosted on S3 by Ashlesh Pandey</h1>
6   </body>
7 </html>
```

Listing 1: HTML code for the static website

Hosting a static website isn't however limited to a single HTML file, rather different styling codes in CSS and functional codes in JS codes are necessary. It should be noted that the entire project must be made public so that the static website functions properly. Similarly, if there are any media assets that are used in the website, they too need to be accessible publicly. The steps to host a static website on AWS S3 include,

1. Create a bucket and make sure the bucket can have public objects.
2. Upload code files to the bucket and make the files public.
3. Access the website using the S3 object url.

Step 1: Log in to AWS console

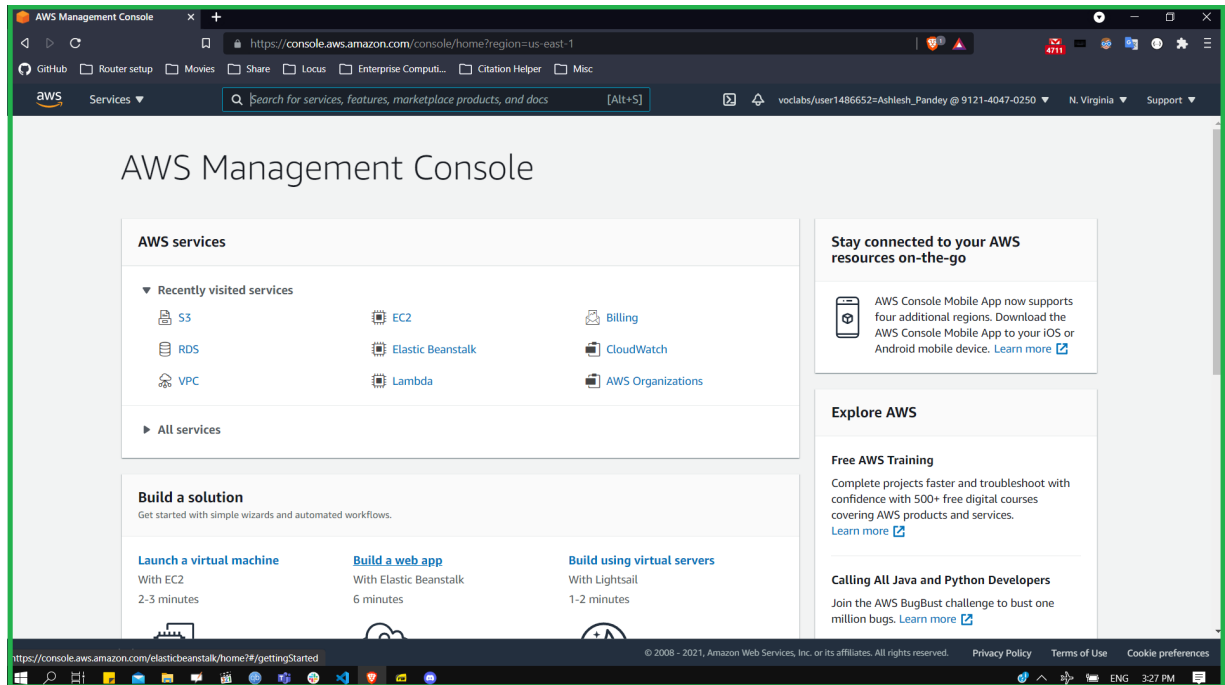


Figure 8: Log in to AWS console

Firstly, log in to the AWS console using a valid user with access to S3 services. In this report, the log in was done using the sandbox account from AWS Academy.

Step 2: Open S3 service from the AWS services list

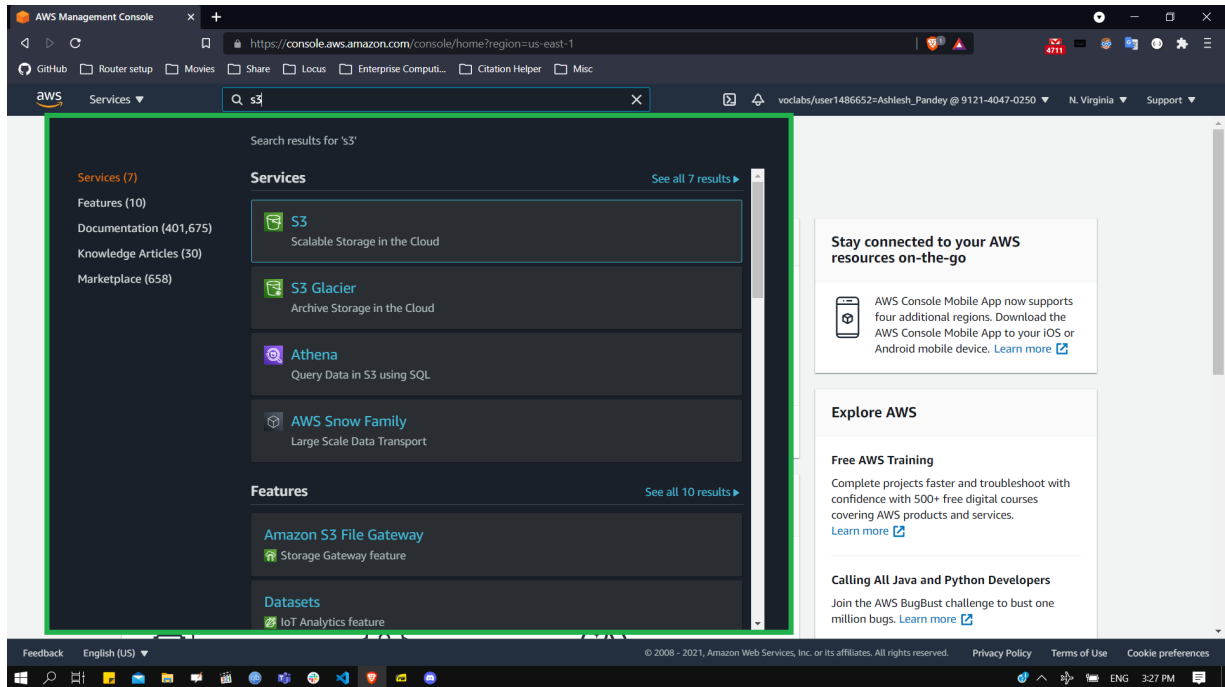


Figure 9: Open S3 service from the AWS services list

Search for S3 service in the search box and select S3 service to proceed.

Step 3: Create a S3 bucket

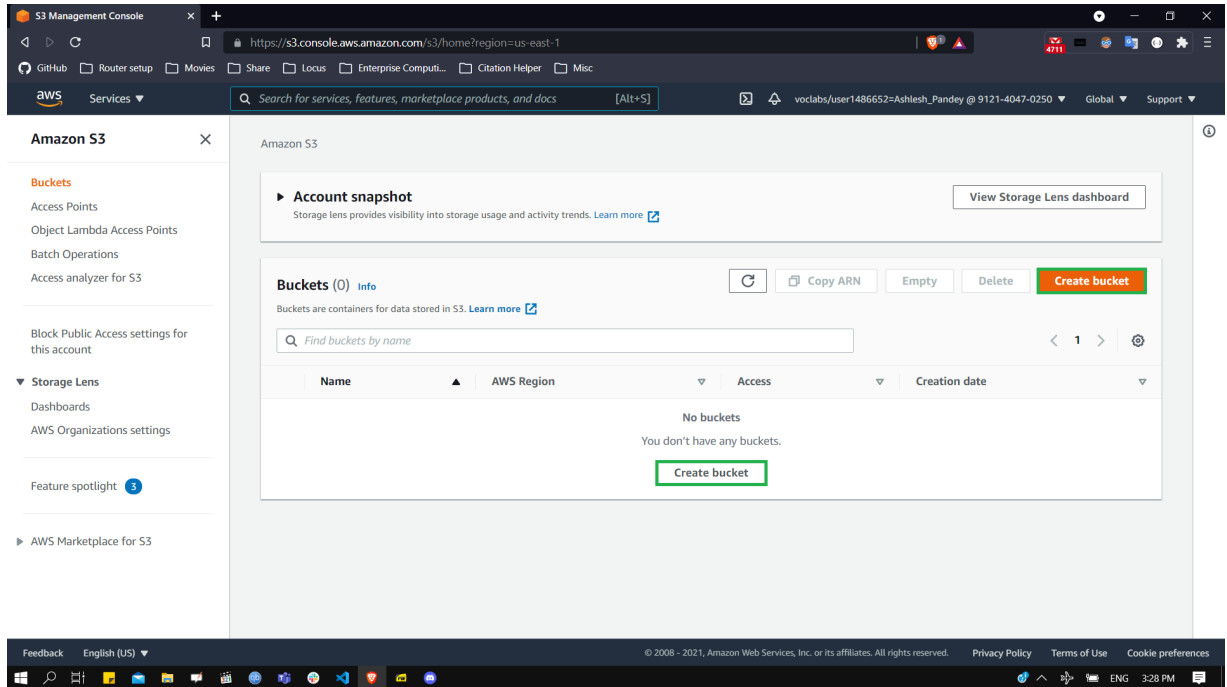


Figure 10: Create a S3 bucket

To create a S3 bucket select any one of the *Create bucket* buttons.

Step 4: Set an unique bucket name

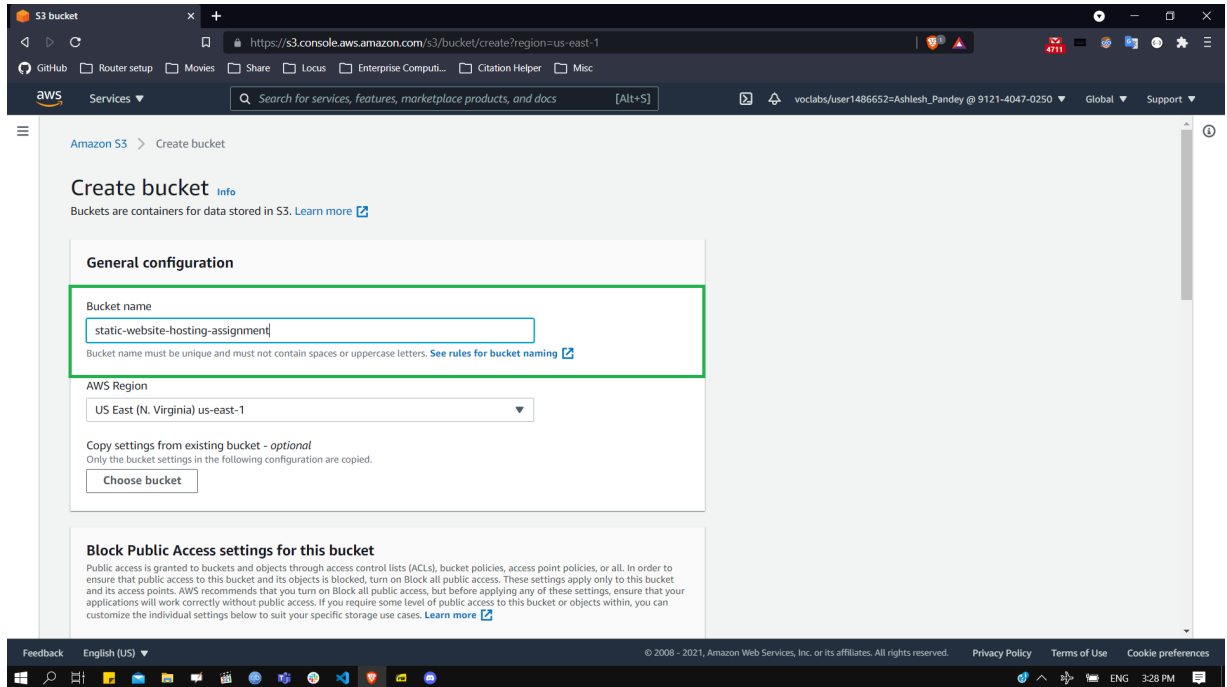


Figure 11: Set an unique bucket name

Bucket names must be unique without using any spaces or uppercase letters. Choose an appropriate name that will be easy to remember. For this example, the name chosen is *static-website-hosting-assignment*.

Step 5: *Grant public access to bucket*

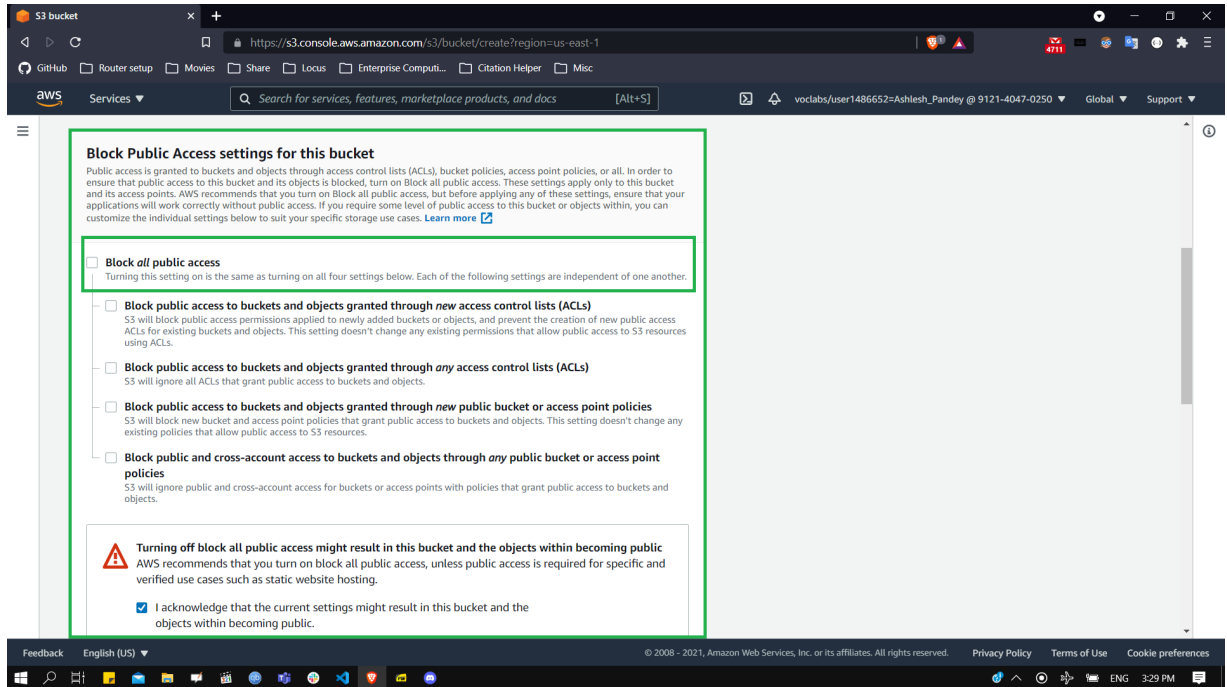


Figure 12: Grant public access to bucket

Unselect the *Block all public access* option. Also select the *I acknowledge that the current settings might result in this bucket and the objects within becoming public* option.

Step 6: Complete bucket creation

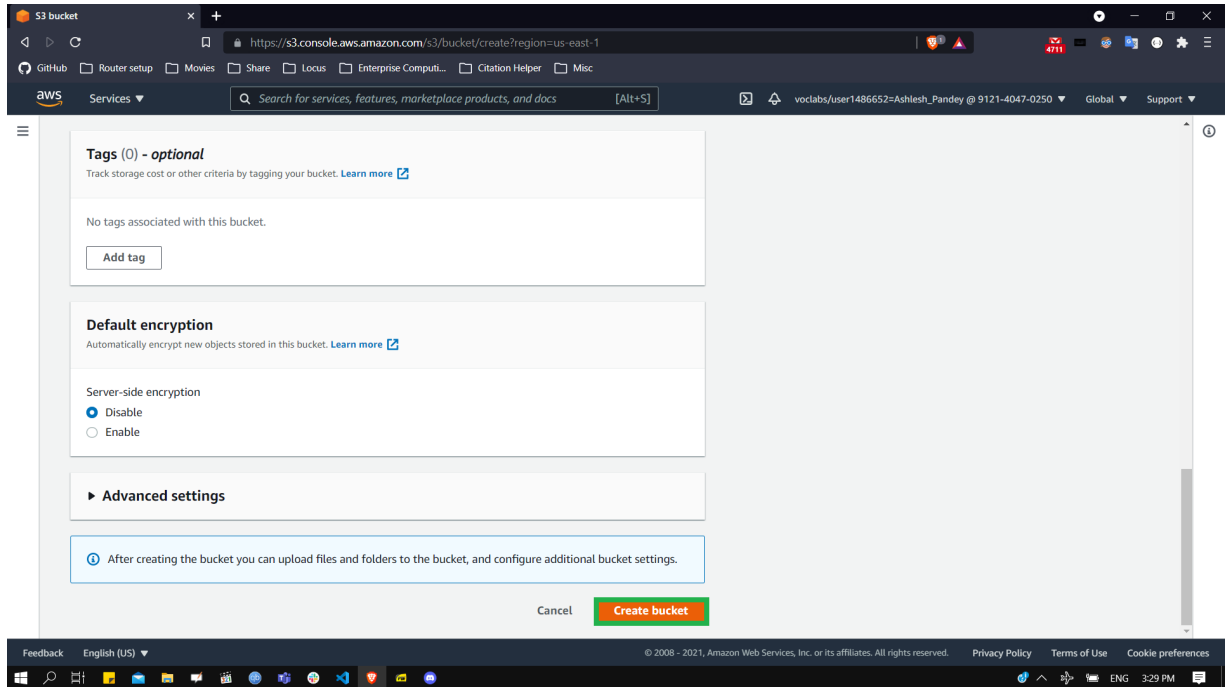


Figure 13: Complete bucket creation

Click on the *Create bucket* button to complete the bucket creation process.

Step 7: *Select the created bucket*

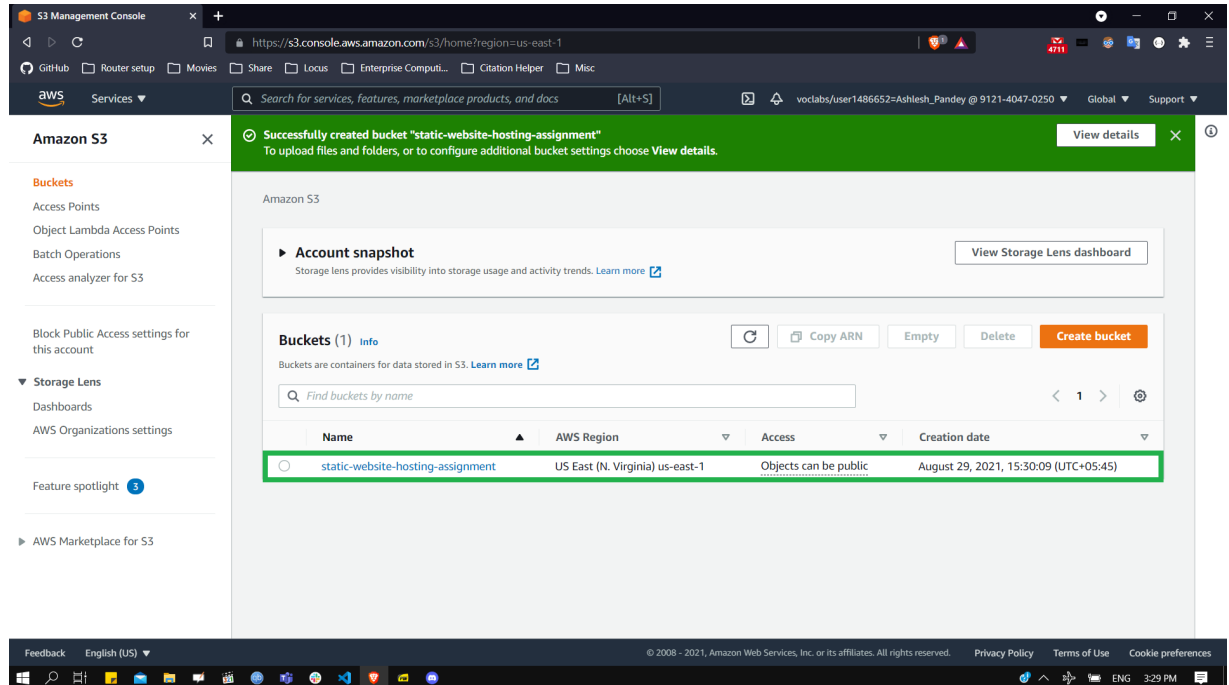


Figure 14: Select the created bucket

Confirm that the bucket is created and then select the bucket to proceed.

Step 8: Upload objects to the bucket

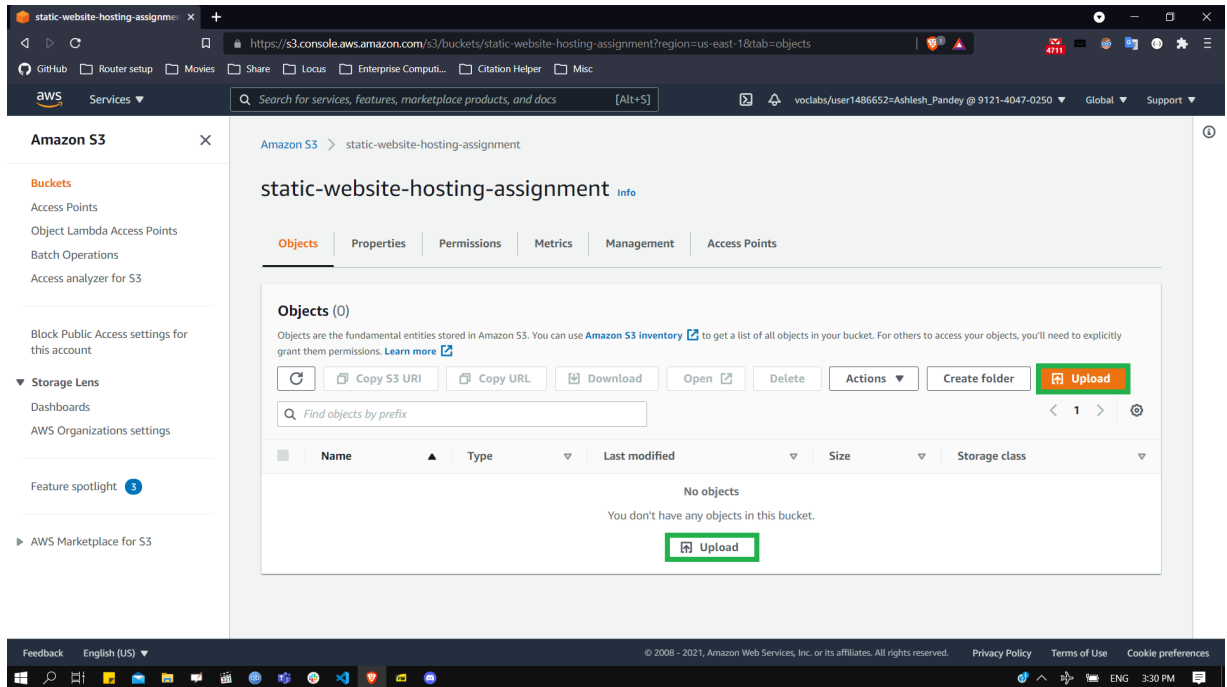


Figure 15: Upload objects to the bucket

Click on the *Upload* button to start the upload process.

Step 9: Upload a file or a folder as needed

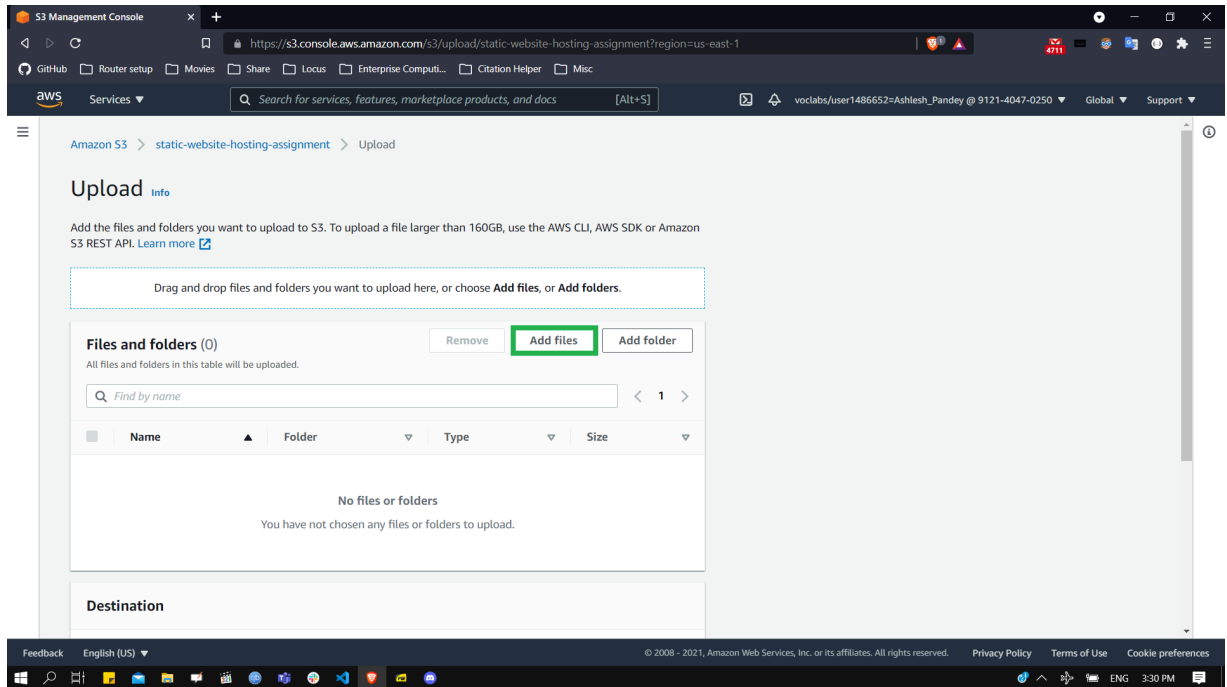


Figure 16: Upload a file or a folder as needed

Click on the *Add files* or *Add folder* button as necessary. Since there is only one file for this example, *Add files* is selected.

Step 10: *Select the file to upload*

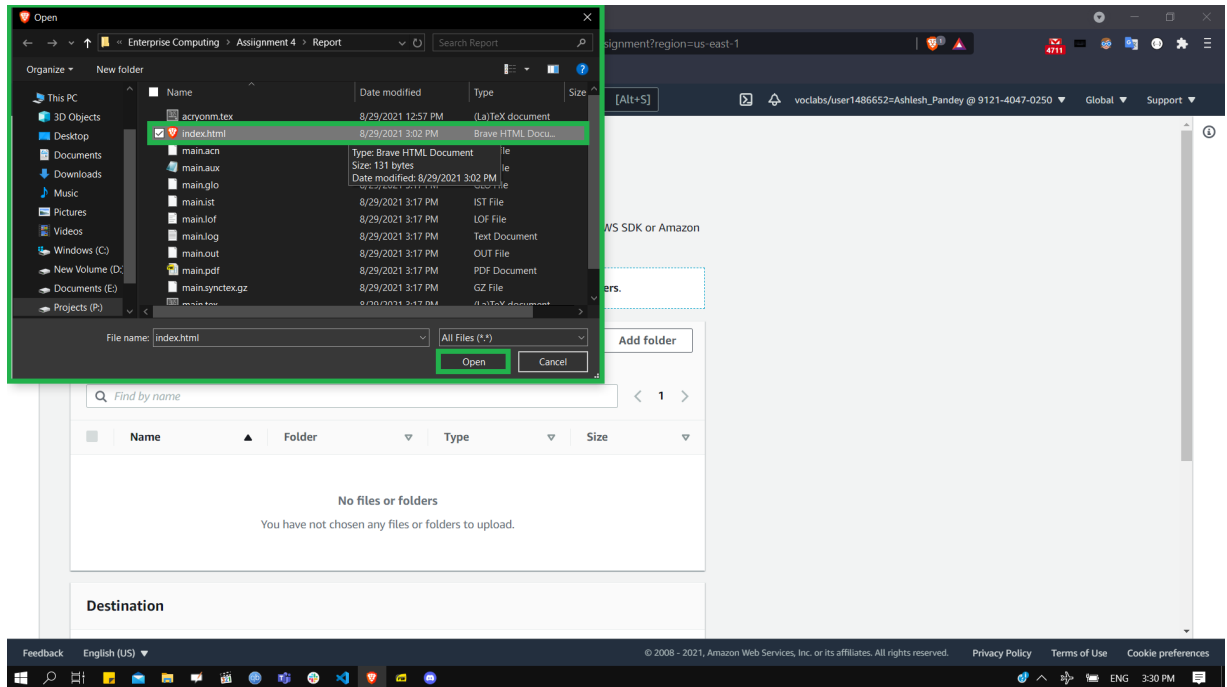


Figure 17: Select the file to upload

Using the prompt that pops up, browse the local directory and select the required file to upload.

Step 11: Upload selected file

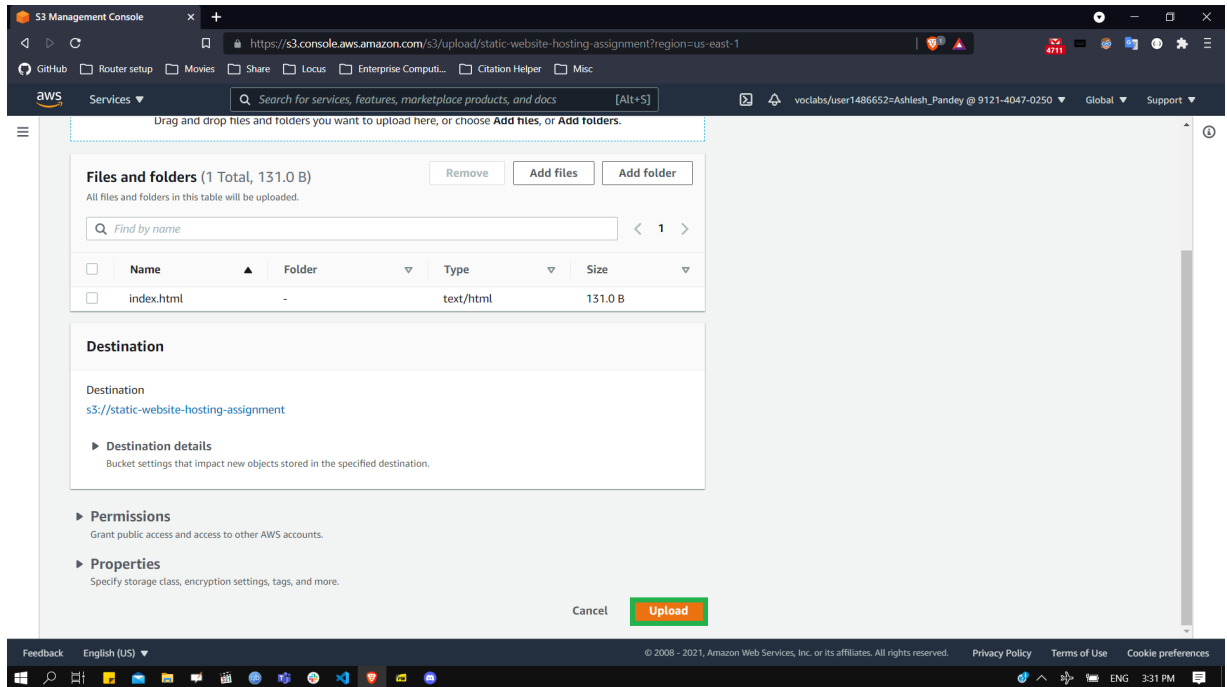


Figure 18: Upload selected file

Click *Upload* button to start uploading the selected file.

Step 12: *Confirm the file upload*

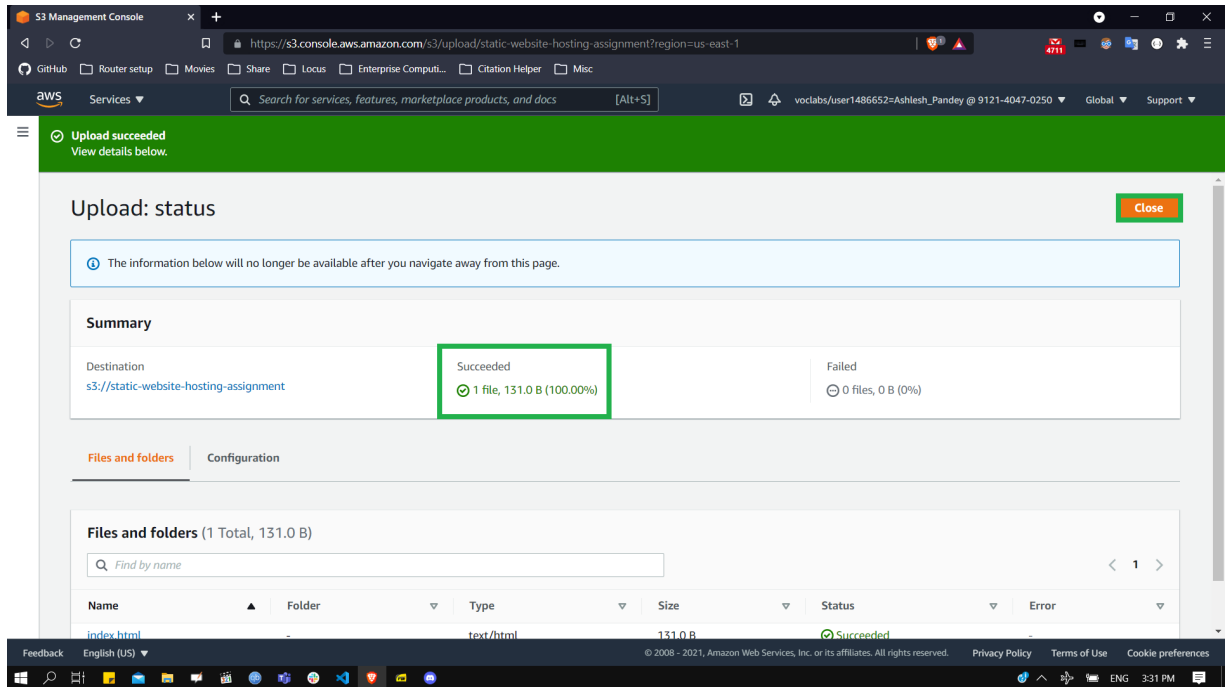


Figure 19: Confirm the file upload

Confirm that the file upload succeeded and then click *Close* button.

Step 13: *Make the object public*

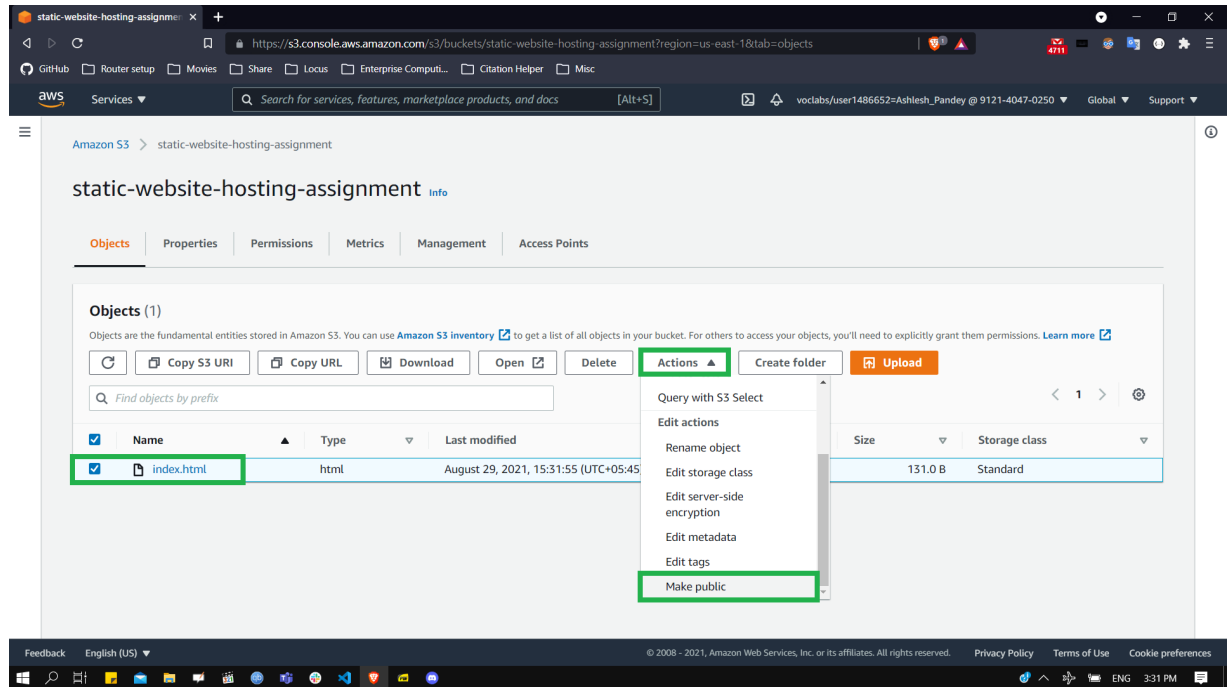


Figure 20: Make the object public

Select the uploaded file, and click *Actions* to open a drop-down menu. Scroll to the bottom of the menu and select *Make public* option.

Step 14: Review the selected object

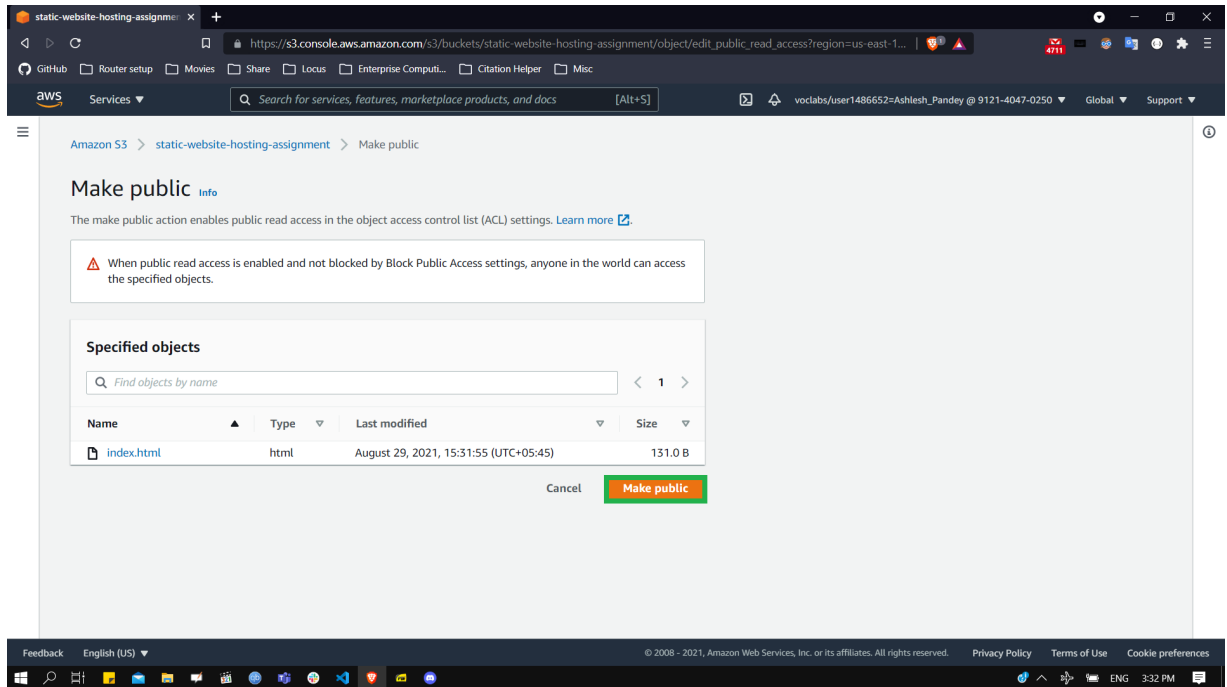


Figure 21: Review the selected object

Review the selected object, and click *Make public* button to proceed.

Step 15: *Confirm public access*

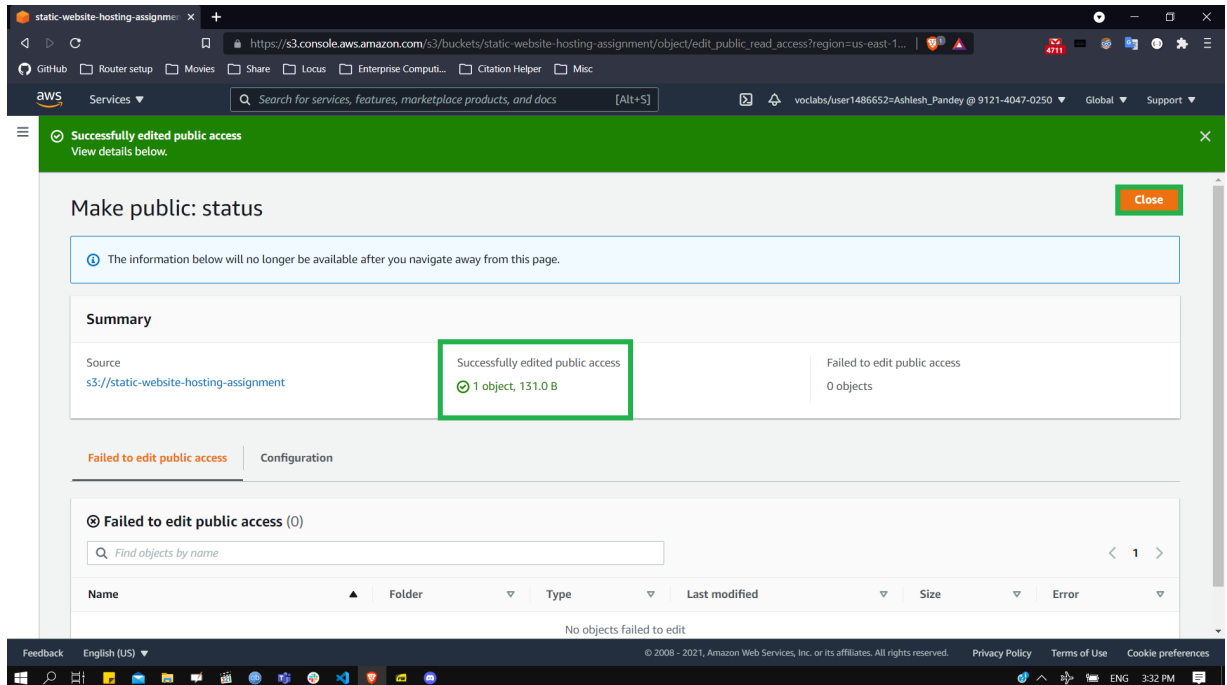


Figure 22: Confirm public access

Confirm that the public access status is edited correctly and then click *Close* button.

Step 16: *Copy the object url*

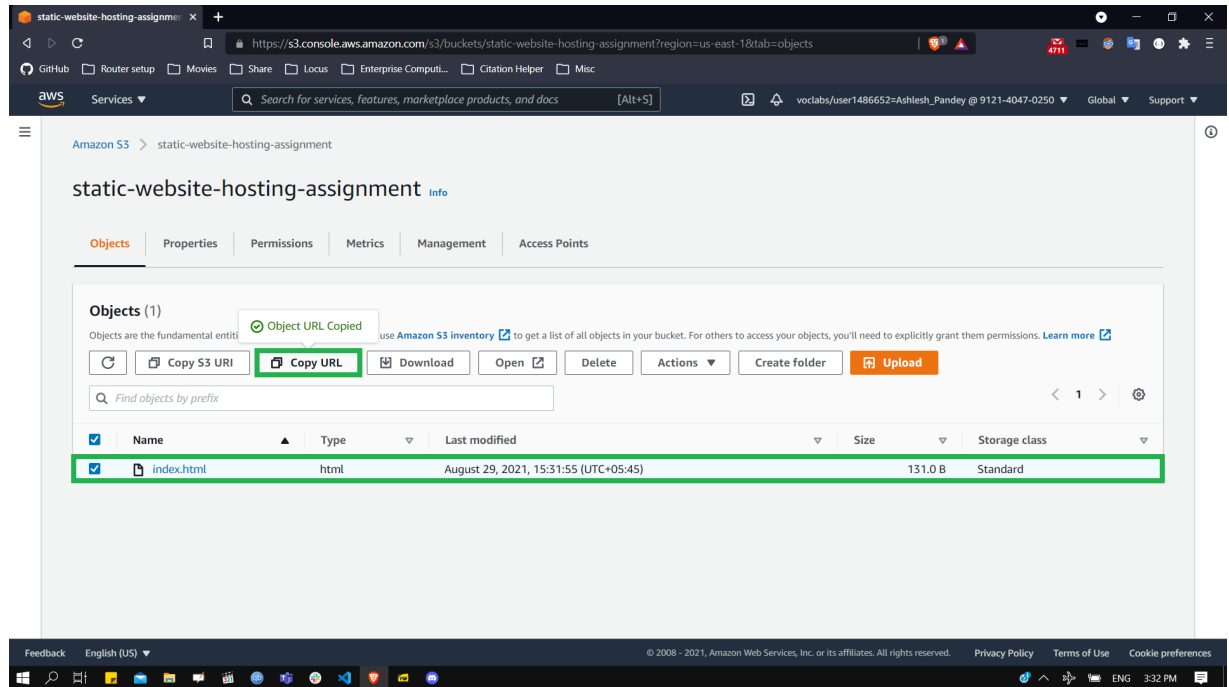


Figure 23: Copy the object url

Select the object and click *Copy URL* button.

Step 17: *Visit the url and access the website*

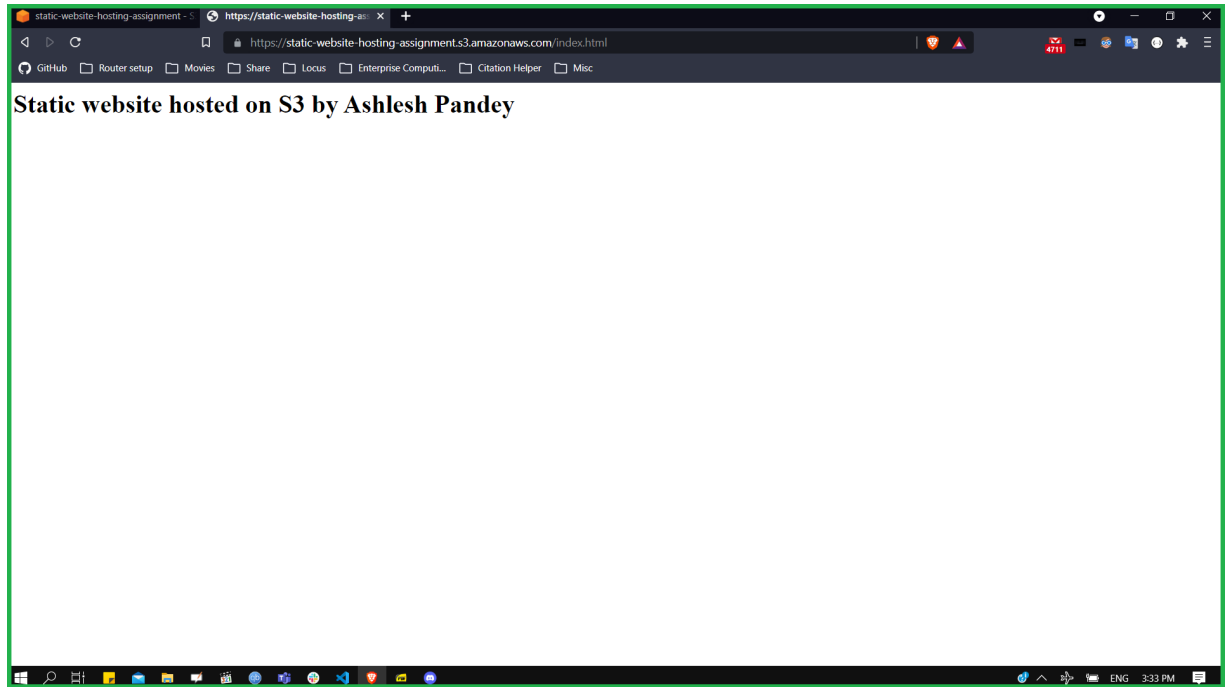


Figure 24: Visit the url and access the website

Paste the copied url in the browser search bar and go to the url. The static website is accessible using this url over the internet.